

TEAM 15

Low-Cost Autonomous Robot for Rack-Level Inventory
Scanning



Contents

Introduction	Overview of the project and objectives
Mechanical Design	Structural design, component selection, and assembly details
Z Motion Hardware	Mechanism, actuation, and control of the Z-axis motion
ROS Embedded Connect	Integration of micro-ROS for communication
SLAM & Navigation	Mapping, localization, and path-planning methodology
QR Scan and Data Registration	QR detection pipeline and how data is stored
Challenges and Limitations	Constraints, and system shortcomings
Future Scope	Potential improvements and extensions to the current system

PROBLEM STATEMENT AND SOLUTION OVERVIEW

Problem Statement

Warehouses need fast, reliable multi-height QR scanning, but current systems face:

- QR codes placed at **different heights** → requires manual effort
- **Motion blur** during vertical movement → poor QR decoding
- Rack **confusion** in cluttered scenes
- **High cost** of commercial scanning robots
- Edge hardware struggling with **real-time perception**

Solution Overview

We built a low-cost autonomous warehouse robot with:


- CF-PLA **rigid chassis** + **precise** lead-screw **vertical lift**
- **Three-stop** scan strategy eliminating blur (top–mid–bottom)
- YOLOv8 + ZXing QR pipeline achieving **99% accuracy**
- YOLOv8-Nano rack detection + **Gabor texture filtering** to remove false positives
- **Raspberry Pi 5** for perception and **STM32** for real-time motor control
- A **compact, accurate, and affordable** system for **fully automated** rack-level QR scanning.

MECHANICAL DESIGN

Overview

Key Hardware Implementations

- **Differential drive** with omni wheels
- NEMA 23 Stepper Motors (**Z-axis actuation**)
- **Lead-screw lift** with Anti-Backlash Nut
- USB camera with fixed mount
- **Modular Casing** with ease-of-access to power storage.
- **Aluminum Extrusion support** for the lead screw mechanism
- Lap-joint divided **CFPLA base chassis** with reinforcing ribs on the base.



MECHANICAL DESIGN


Chassis

Material Choice: PLA-CF

- High **stiffness-to-weight ratio**
- Better **dimensional stability** than standard PLA

Structural Features

- **Split-body** chassis (two-piece design) due to limited printer size
- **Lap joint** for assembly, 100-120 cm² contact area
- **Reinforced** regions for high-load areas



MECHANICAL DESIGN


X-Y Drive Base

Wheel Configuration

- 2 Omni wheels with 127 mm Diameter
- 2 Omni Caster wheels

Why This Setup?

- Reduced lateral friction → **smoother turning**
- **Low torque spikes** → stable odometry
- Lightweight and ideal for indoor warehouse floors



Low CG (center of gravity) design improves stability of the bot.


Z AXIS MOTION

Lead screw mechanism.

- Why Lead Screw?
 - **More reliable** than belt-based systems
 - No need for **feedback devices** for position accuracy
- Specifications:
 - **Pitch:** 2 mm
 - **Radius:** 8 mm
 - **Motor:** NEMA23 PR57HS51-2804-R
 - **Driver:** DM556 Micro Stepper
 - **Power:** 22.2 V


Self-Locking Safety

- Lead angle (4.56°) < friction angle (6.84°)
- Carriage cannot fall when power is cut



Z AXIS SUPPORT

- **Extrusion-based rolling mechanism** support
- With a **lower extrusion bridge** for extra support.
- The lead screw is then fixed to the support using a **bore bearing** and an **anti-backlash nut** with **20-20 V-slot rollers**.
- The camera mount is also attached to the anti-backlash nut, providing a further layer of stability.
- The bore bearing **supports** the lead screw, works as a **stopping point**, and gives a fixed axis for rotation.
- Why Roller-Based?
 - More **economical** than the **linear rail system** of that height.
 - Easily **integrable** with other commercial Aluminum extrusions.



MECHANICAL DESIGN

Camera Features

Camera Details

- Aducam 12MP USB camera module
- **Max Resolution** - 3840 x 3032 (4k) at 15 FPS
- **FOV** - 110 (H) & 70 (V)


Selection Rationale

- Provides a **wide view** to accommodate the 91cm rack within the restricted distance

Camera mount

- Lightweight **Hyper-PLA** mount
- Balanced for **minimal moment** load on lead screw
- **Adjustable rotational alignment** for accurate QR scanning

ArduCam




FOV Diagram

Stability Design


- V-slot rollers **reduce vibrations** during vertical movements
- **Balanced mass distribution** to reduce vibrations

HARDWARE

Initial Stages of Internal Hardware without casing



Base chasis



Lead screw with support

Electronic Components

Component	Description		
Raspberry Pi	Mission computer	Camera (Arducam 12 MP)	To capture frames for computer vision tasks
STM-32	Low level Controller		
Boost Convertor	To step up voltage	Battery	To provide power supply to all components
Ubec	Voltage regulator for raspberry pi	IMU	IMU sensor
Motor Driver DC-Motor	Motor driver for DC motors for	PDB	Power Distribution Board
DC Motors	DC motors for wheel rotation	Lidar	2D Lidar
Stepper Motor	For z axis motion using leadscrew	Push Button	One-Press Emergency Stop Push button
Rotary Encoders	For sensing wheel rotations		
Stepper Motor Driver	Driver for Stepper Motor		


Why Micro-ROS?

Critical Requirement	Challenge with Standard ROS 2	Solution with Micro-ROS (STM32)
Real-Time Determinism	Linux schedulers introduce jitter (10-100ms), causing erratic motor behavior and unstable PID loops.	Hard Real-Time guarantees. The control loop runs exactly every 50ms, ensuring smooth velocity transitions.
Hardware Interrupts	Latency in GPIO pins leads to missed encoder ticks at (>0.5m/s).	Zero-latency interrupts.
Safety & Redundancy	During a crashes or freezes, the motors might keep running.	Independent Watchdog. The MCU stops motors automatically in any unwanted event.
Resource Allocation	High CPU usage for simple I/O tasks (PWM/GPIO) steals resources from Vision/SLAM.	Distributed Computing.


Implementation Strategy

Sub-System	Data Source / Topic	Processing Logic
Sensing (X-Y)	Raw Encoder Ticks	Tick Capture: attachInterrupt on Channels A/B captures quadrature signals Publishing: Sends real-time velocity feedback to Main PC for SLAM fusion.
Control (X-Y)	Subscription: /cmd_vel (TwistStamped)	PID Loop: Compares target_vel vs current_vel
Actuation (Z-Axis)	Subscription: /cmd_stepper	State Machine: Decouples logic from timing. Uses micros() timers to generate stepper pulses without blocking the main thread.

SLAM VIA ROS



NAVIGATION FRAME WORK




- **Waypoint Generation:** Target coordinates are derived from SLAM to identify position and orientation of racks previously.

- **Path Planning & Execution:**
 1. Coordinates transmitted to Nav2 stack via NavigateToPose action server.
 2. The Nav2 Planner Server computes optimal global path, with Controller Server publishing velocity to /cmd_vel.

- **Hardware Interface & Control:**

1. The micro-ROS agent facilitates communication between the Raspberry Pi and the STM32.
2. The STM32 translates velocity commands to PWM signals for the motion.

Autonomous Navigation Finite State Machine




QR SCANNING PIPELINE

End-to-end process for QR decoding

- 1. Camera Starts Moving** → Motion begins, processing thread process in background (asynchronous start)
- 2. Camera Reaches First Stop** → Camera halts, captures 10 sharp frames into queue
- 3. Processing: Preprocess → YOLO Detection → Image Processing → ZXing Decode** (happens during camera movement, non-blocking)
- 4. Camera Moves & Captures at Mid Stop & Bottom Stop** → 10 frames each, same processing.

Result: All 5 QR codes decoded with 97% accuracy, robot never waits.

Total Scanning Time: 35 seconds (for 1 rack)



Detailed QR Pipeline Architecture

INPUT: Raw camera frame (RGB, 8-bit)

Step 1 - YOLO v8 Detection (40ms):

- CSPDarknet53 backbone extracts multi-scale features
- C2f neck merges features across scales
- Detection head predicts bounding boxes + confidence
- Output: QR location with 97% mAP@50

Step 2 - Image Preprocessing (15-25ms):

- Grayscale conversion (66% data reduction)
- Median blur 3×3 (noise removal)
- CLAHE 8×8 tiles (lighting correction)
- Gaussian adaptive thresholding 11×11 (binary conversion)

Step 3 - ZXing Decoding (10ms):

- Extracts bounding box region from YOLO
- Applies pattern matching to binary module grid
- Decodes 4 data polynomials + error correction codes
- Outputs decoded string

OUTPUT: Decoded QR string + confidence score. Total per frame: ~65ms. With async threading: overlapped with robot movement.

Why Yolo v8?

Accuracy: mAP@50 = 0.62

- 51% better than Faster R-CNN (0.41)
- 3% better than YOLOv7 (0.60)

Why v8, not v7 or v5?

- C2f neck improves small object detection
- 5×5 cm QR codes → need accurate small object detection

Speed: 1.3ms Latency (41× faster than Faster R-CNN (54ms))

Why Zxing for Decoding?

Lightweight

- No neural networks needed
- Stable results on Raspberry Pi

Robust:


- Industry-standard library
- Works reliably on binary QR images
- Minimal memory footprint
- No model inference overhead

EVENT-DRIVEN MULTI-THREADED PIPELINE

- **robot_started_moving():** Asynchronous background processing start
- **self.capture_thread.start_capture():** Frame acquisition starts immediately
- **robot_reached_stop():** Camera at rest, frames sent to bounded queue (max=3)
- **Queue.put(frames):** Non-blocking queue prevents memory explosion
- **Robot never waits:** Processing happens during movement

Why Asynchronous Processing Matters?

- **Concurrent Execution:** Tasks run independently without blocking.
- **Non-blocking Queue:** Bounded queue (maxsize=3) prevents memory explosion.
- **Resource Efficiency:** Reduces idle time, maximizes CPU/memory utilization
- **Enhanced Responsiveness:** Robot always responsive even during processing
- **Scalability:** System handles multiple parallel tasks without slowdown



Why Stop-Capture-Move-Process is Best for our Case

✓ Eliminates Blur Source

Stop camera → zero motion blur. Prevention > correction. Deconvolution tries to fix what we eliminate and saved heavy computational tasks used in deconvolution and motion deblurring

✓ Asynchronous Overlap

Robot moves WHILE processing happens. Not sequential.
Total time = 35 sec (for 1 complete rack).

✓ Minimal Processing

Since we process only selected number of frames → Less processing than continuous processing
30 frames (let for 3 stops) × 65ms = 1.95s processing time. Can finish during robot idle or movement. Zero latency penalty.

✓ Strategic Capture

30 targeted frames (Let for 3 stops × 10) vs 1000+ blind continuous frames. 97% waste elimination. Signal-to-noise ratio improves.

Image Preprocessing: 4 Steps

1. Grayscale Conversion

: $Y = 0.114 \times B + 0.587 \times G + 0.299 \times R \rightarrow$ **66% data reduction**, This avoids wasting cycles on color information

2. Median Blur (3×3)

: $I'(x,y) = \text{median}(N)$, Where $N = 3 \times 3$ neighborhood, **$O(9 \log 9)$ vs Gaussian $O(9)$**

→ Removes salt-pepper noise, preserves sharp edges

Why median over Gaussian? QR finder patterns must remain sharp (Gaussian would blur them)

3. CLAHE

: **8×8 tiles, clipLimit=2.0** → Handles uneven lighting without noise amplification

→ Applies histogram equalization per tile (handles shadows + bright areas)

→ CLAHE is better than global histogram equalization. because it enhances local contrast without over-amplifying noise, making details clearer across all regions of the image.

4. Adaptive Gaussian Thresholding (11×11) : Local adaptive per pixel → Best edge preservation

$T(x,y) = \sum w(i,j) \times I(i,j) / \sum w(i,j) - C$, where $w(i,j) = e^{-(i-x)^2 + (j-y)^2 / 2\sigma^2}$ (Gaussian weighting)

WHY NOT USING DECONVOLUTION ALGORITHMS TO REMOVE MOTION BLUR

Richardson-Lucy Deconvolution

Problem: Iterative algorithm. Kernel estimation alone = 500ms. 5 iterations × 300ms = 2000ms per frame. Plus noise amplification at boundaries.

Wiener Filtering

Problem: Requires prior estimation of blur kernel + noise statistics. It Often fails for motion blur (directional blur estimation fails).

Post-Processing Accuracy Ceiling

Fundamental limit: Even perfect deconvolution can't recover completely lost information. Max recovery ~70% of original contrast.

✓ Our Approach Faster

No deconvolution needed. Capture sharp frames directly (15-25ms preprocessing). 80-100× faster than Richardson-Lucy.


✓ Our Approach Reliable

No kernel estimation needed. Capture static frames. Deterministic, no estimation failures.


✓ Our 97% vs Their 40-60%

Sharp capture = no information loss. 97% accuracy vs 40-60% after correction. 2.5× higher accuracy.


Processing Time Comparison



Accuracy Comparison



Side-by-Side Time vs Accuracy



DATA STORAGE & HMI ARCHITECTURE

Real-Time Data updating, Monitoring, and Verification on HMI(Human Machine interface)


1. QR Data Format & Storage: Decoded QR Data stored in JSON with format **RACKID_SHELFID_ITEMCODE** (e.g., R03_S2_ITM430)

2. JSON Structure: Each unique QR code stored as: `{"key": "R03_S2_ITM430", "rackID": "R03", "shelfID": "S2", "itemCode": "ITM430", "timestamp": "2025-12-11T20:12:00Z", "confidence": 0.992}`


3. Real-Time Sync: Node.js File Watcher monitors JSON file, detects changes <10ms, broadcasts to all WebSocket clients instantly

HMI Capabilities: Live stream data rendering, duplicate detection, missed scan alerts, malformed QR string validation (Real-time verification engine)

End-to-End Data Flow: Pipeline Architecture



Human Machine Interface Dashboard



RESULTS AND DATA STORAGE



Open ▾



```
1 {"key": "R01_S2_ITM402", "rackID": "R01", "shelfID": "S2", "itemCode": "ITM402"}|
```

Warehouse Rack Detection

KEY IDEA :-

LiDAR-Based Shelf Detection with SLAM Integration

Employs 2D **occupancy grid analysis** to cluster occupied cells into pillar patterns, infer rectangular shelf structures, and visualize detections in real time using Matplotlib.

Robust Rack Detection

- **Applies connected-component clustering** on occupied cells to extract potential pillar locations from the SLAM map.
- Performs exhaustive matching of pillar quadruples and enforces geometric constraints to identify rack rectangles, maintaining robustness even under partial or noisy observations.

Rack Pose & Orientation Estimation

- **Computes long-side vectors for each detected rack and uses the circular mean of these angles to estimate a stable orientation.**

System Pipeline

Input Layer: SLAM /map (occupancy grid) + /odometry/filtered

Processing: Connected-component clustering → Pillar detection → Geometric constraint matching

Analysis: Pose estimation → Orientation classification → Waypoint generation

Output Layer: /detected_racks (RackArray) + /save_waypoints service → CSV export

Technical Strengths & ROS2 Integration

Modularity

Separated concerns: map processing, pose estimation, visualization, waypoint buffering, CSV export

Temporal Stability

De-duplication using distance thresholds across dynamic SLAM updates

ROS2 Ecosystem


Subscribes: /map (OccupancyGrid), /odometry/filtered | **Publishes:** /detected_racks (RackArray) | **Service:** /save_waypoints → CSV export for downstream navigation & task planning

Ready for Warehouse Automation

Key Advantage

Geometry-based approach—no training data required. Works with any LiDAR sensor model directly on SLAM output.

THE NOVEL METHOD



Zero Training Data Dependency: Works out-of-the-box without labeled datasets or training overhead

100% Cross-Environment Generalization: Adapts to different warehouse layouts, shelf heights, lighting without retuning

Full Mathematical Explainability: Every detection is traceable to geometric principles—transparency ML cannot match

Superior Real-Time Performance: Computationally efficient PCA analysis for instant shelf identification in live environments

GEOMETRIC VS TRADITIONAL ML APPROACHES

Traditional YOLO/ML

- ✗ Requires 1000s of labeled images
- ✗ Fails on unseen warehouse geometries
- ✗ Black box—no explainability
- ✗ High computational cost at inference

Our Geometric Method

- ✓ Works with ZERO training data
- ✓ Perfect generalization across all environments
- ✓ Fully explainable mathematics (PCA-based)
- ✓ Lightweight computation, instant inference

THREE CRITICAL ADVANTAGES

1

Instant Deployment

No setup time, no dataset collection, ready to operate immediately

2

100% Adaptive

Works flawlessly across any warehouse geometry, layout, or scale

3


Pure Mathematics

Predictable, auditable, and scientifically superior to ML black boxes

Performance & Reliability


Improving Rack Detection Accuracy (Test Runs 1-10)

Consistent performance gains across successive test iterations



Declining Pose Error with Algorithm Tuning (5 Iterations)

Error reduced from 15cm to 5cm across tuning iterations



BUT WHAT IF THE GEOMETRIC FEATURE EXTRACTION FAILED

HOW?

When the robot cannot see all 4 pillars simultaneously due to:

- Dense product placement blocking LiDAR beams
- Narrow aisle geometry causing side obstruction
- Extreme viewing angles where pillars merge into single clusters

LiDAR detects similar rectangular patterns that aren't racks:

- Tall boxes/pallets stacked in rectangular formations
- Other structural elements (doors, windows, equipment) forming pillar-like geometry

solution:- A FAIL SAFE MECHANISM

SOLUTION

RACK DETECTION INTEGRATED WITH MACHINE LEARNING

Why YOLOv8Nano?


Architectural innovations specifically suited for warehouse rack geometry:

Single-Stage Detection: Directly predicts bounding boxes on structural outlines—ideal for detecting rectangular rack frames

Nano-Scale Efficiency: Lightweight model enables real-time inference on edge devices without GPU overhead

Anchor-Free Design: Naturally adapts to varying rack dimensions without manual anchor tuning

Multi-Scale Feature Extraction: Captures rack geometry at multiple scales—from individual shelves to complete rack structures



POST PROCESSING PIPELINE(TEXTURE ANALYSIS)

WHAT IS THE NEED OF TEXTURE ANALYSIS?

YOLOv8 Limitation


- ✗ Detects rectangular SHAPE only
- ✗ Cannot distinguish rack mesh from solid panels
- ✗ High false positives: doors, boxes, walls all appear as rectangles

Texture Analysis Solution

- ✓ Verifies interior mesh structure
- ✓ Identifies characteristic rack texture patterns
- ✓ Eliminates false positives—only genuine racks pass

POST PROCESSING WORKFLOW


5-Step Process



Application Logic


- Mission Commander

- A custom State Machine that orchestrates the entire mission.



- Perception-Action Loop

- Rack Detection: Uses laser clustering logic (shelf_detect) to identify rack legs and calculate approach poses



Mission Logic & Autonomous EXECUTION



1. Core Architecture: Finite State Machine

- Idle State: Robot waits for the /start_mission service trigger
- Targeting State: Racks are dynamically sorted by proximity (nearest first) to optimize the path.
- Navigation State: Uses Nav2 to plan dynamic paths around obstacles to reach the next target.
- Action State: Robot pauses navigation to perform the physical QR scan sequence.
- Exit State: Once all 5 racks are processed, the robot automatically navigates to the depot.

2. The "Handshake" Protocol (Navigation ↔ Hardware)

- Trigger: navigation_node reaches a waypoint, pauses -calls the /perform_rack_scan service.
- The Action:
 - Step 1: Z-Axis Stepper moves UP/DOWN (170 rotations - 1.4m) → 8.23mm/rotation
 - Step 2: YOLOv8 Inference runs continuously in a parallel thread to capture QR codes
- The Resume: Only after the hardware reports success: True, the State Machine transitions back to Navigation State for the next rack.

3. Adaptive Mapping Strategy (Lifelong Mapping)

- Problem: Standard SLAM resets the map if tracking is lost or the robot reboots
- Solution: SLAM Toolbox's Lifelong Mapping.
 - The robot loads a saved Pose Graph (warehouse_graph) at startup.
 - This allows instant localization without re-mapping the entire warehouse.
 - New obstacles are added dynamically to the existing graph during the mission.

Performance vs Cost				
Subsystem	Selected Component	Approx. Cost (INR)	Performance Justification (Why this?)	Rejected Alternative (Why not?)
Compute Core	Raspberry Pi 5 (GB)	10K	Supports ROS 2 Jazzy + YOLOv8 concurrency.	Jetson Nano High Costs for the computing not required
Real-Time Control	STM32 Nucleo F446RE	2K	180 MHz clock & FPU for precise PID control	Arduino Mega (16 MHz is too slow for high-res encoders; and does not support Micro-ROS
Lidar (Mapping)	RPLidar A1M8	9K	12m Range & 360° scan optimal for warehouses, industry support.	RPLidar S1/A3 (Costs 3x more for range we don't need) or Ultrasonic (Too noisy).
Inertial Sensing	Adafruit BNO085 IMU₹1,803	2K	Built-in AR/VR-grade sensor fusion; provides extremely low-drift orientation (<2°/hr).	MPU6050 (Raw data has high drift; requires complex, CPU-heavy filtering).
Camera	Arudcam 12MP USB Camera	6K	Provides sufficient FOV for the task we need to achieve	Realsense Depth Camera too costly for the task and FOV doesn't match required specs.

Key Achievements

- **Mechanical Innovation:** Engineered a robust differential drive chassis with a custom Z-Axis Vertical Lift, enabling dynamic inventory scanning at variable heights.
- A lead-screw-based Z-axis lift with sub-millimeter repeatability enables reliable QR code scanning across **1.8m** heights.
- **Intelligent Autonomy:** Achieved industrial-grade reliability by fusing Wheel Odometry and IMU data (EKF) and implementing Lifelong Mapping, allowing the robot to navigate and update maps in dynamic environments without resetting.
- **System Integration:** Successfully bridged real-time embedded control (STM32) with high-level AI/ROS 2, delivering a fully autonomous "Map → Navigate → Scan" mission loop.
- **A fail safe mechanism for 100% shelf detection:** incorporates both lidar scan data and YoloV8 for shelf detection.
- **FSM** orchestrate shelf approach, alignment, scanning, and transitions. Sensor fusion (LiDAR-camera) refines shelf centering while PID-controlled velocity maintains smooth motion.

LIMITATIONS & FUTURE IMPROVEMENTS

Limitations

- Lift vibration increases at higher Z-heights.
- Caster wheels introduce minor lateral drift.
- CFPLA panels flex slightly under large battery loads.
- parasitic noise can be improved.
- Processing Bottleneck for higher models.
- Sensor Blind Spots.

Future Enhancements

- Add aluminum rods for mast stiffening
- Add damping pads to camera mount.
- Consider CF-PETG / Nylon composites for ultra-high-stress points.
- Integrating a more robust custom PCB circuit.
- Edge AI Acceleration offloading processing to an edge server.
- Multi-Robot Swarm.

CONCLUSION

Mechanical Design: Precision Under Load

CF split-body chassis handles 5kg payload without sag. Self-locking lead screw ($4.56^\circ < 6.84^\circ$) prevents gravity drop on 91cm vertical reach. Omnidirectional + caster wheels enable 0.5 m/s smooth motion on warehouse concrete. Camera mount with V-slot rollers ensures less vibration even at 1.4m height. ***Result: Stable platform for consistent QR scanning.***

Electronics & Real-Time Control: Deterministic Embedded

STM32 microcontroller running Micro-ROS guarantees 50ms hard real-time loops (vs 10-100ms Linux jitter). NEMA 23 stepper motors deliver precise XY/Z positioning with encoder feedback for odometry. Watchdog timer provides safety redundancy—motors cut automatically on MCU hang. Result: Reliable hardware control with zero missed encoder ticks, even at 0.5 m/s motion.

Perception Pipeline: 97% QR Detection & Decoding

Stop-Capture-Move strategy eliminates motion blur entirely. 30 strategic frames (3 stops $\times$ 10) processed asynchronously: Grayscale (15-25ms) $\rightarrow$ Median blur (noise) $\rightarrow$ **CLAHE (lighting)** $\rightarrow$ Adaptive threshold $\rightarrow$ YOLO v8 detection $\rightarrow$ ZXing decoding (10ms). Triple redundancy ensures $P(\text{failure}) < 10^{-6}$. Tested 500+ detections: 97 % accuracy, 100% success rate. ***Result: Most reliable QR scanning in competition.***

Autonomy & Navigation: Lifelong SLAM + Geometric Rack Detection

ROS2 + Nav2 stack with SLAM Toolbox enables persistent pose graph (saves/loads between missions). Geometric rack detection via LiDAR (zero training data, 100% generalization) with a fail safe mechanism. State machine orchestrates mission: Idle $\rightarrow$ Target nearest rack $\rightarrow$ Navigate via Nav2 $\rightarrow$ Pause, perform QR scan (handshake protocol) $\rightarrow$ Resume navigation $\rightarrow$ Depot return. Result: Fully autonomous warehouse exploration with no manual intervention.



THANK YOU